

```
1 import numpy as np
2 import keras
3 from keras import layers
4 import tensorflow_datasets as tfds
5 import matplotlib.pyplot as plt

1 tfds.disable_progress_bar()
2
3 train_ds, validation_ds, test_ds = tfds.load(
4     "cats_vs_dogs",
5     # Reserve 10% for validation and 10% for test
6     split=["train[:40%]", "train[40%:50%]", "train[50%:60%]"],
7     as_supervised=True, # Include labels
8 )
9
10 print(f"Number of training samples: {train_ds.cardinality()}")
11 print(f"Number of validation samples: {validation_ds.cardinality()}")
12 print(f"Number of test samples: {test_ds.cardinality()}")
```

➡ Downloading and preparing dataset 786.67 MiB (download: 786.67 MiB, generated: 1.04 GiB, total: 1.81 GiB) to /root/tensorflow.
WARNING:absl:1738 images were corrupted and were skipped
Dataset cats_vs_dogs downloaded and prepared to /root/tensorflow_datasets/cats_vs_dogs/4.0.1. Subsequent calls will reuse this.
Number of training samples: 9305
Number of validation samples: 2326
Number of test samples: 2326



```
1 plt.figure(figsize=(10, 10))
2 for i, (image, label) in enumerate(train_ds.take(9)):
3     ax = plt.subplot(3, 3, i + 1)
4     plt.imshow(image)
5     plt.title(int(label))
6     plt.axis("off")
```




```
1 resize_fn = keras.layers.Resizing(150, 150)
2
3 train_ds = train_ds.map(lambda x, y: (resize_fn(x), y))
4 validation_ds = validation_ds.map(lambda x, y: (resize_fn(x), y))
5 test_ds = test_ds.map(lambda x, y: (resize_fn(x), y))
```



```
1 augmentation_layers = [  
2     layers.RandomFlip("horizontal"),  
3     layers.RandomRotation(0.1),  
4 ]  
5  
6  
7 def data_augmentation(x):  
8     for layer in augmentation_layers:  
9         x = layer(x)  
10    return x  
11  
12  
13 train_ds = train_ds.map(lambda x, y: (data_augmentation(x), y))
```

```
1 from tensorflow import data as tf_data  
2  
3 batch_size = 64  
4  
5 train_ds = train_ds.batch(batch_size).prefetch(tf_data.AUTOTUNE).cache()  
6 validation_ds = validation_ds.batch(batch_size).prefetch(tf_data.AUTOTUNE).cache()  
7 test_ds = test_ds.batch(batch_size).prefetch(tf_data.AUTOTUNE).cache()
```

```
1 for images, labels in train_ds.take(1):  
2     plt.figure(figsize=(10, 10))  
3     first_image = images[0]  
4     for i in range(9):  
5         ax = plt.subplot(3, 3, i + 1)  
6         augmented_image = data_augmentation(np.expand_dims(first_image, 0))  
7         plt.imshow(np.array(augmented_image[0]).astype("int32"))  
8         plt.title(int(labels[0]))  
9         plt.axis("off")
```



```
1 base_model = keras.applications.Xception(  
2     weights="imagenet", # Load weights pre-trained on ImageNet.  
3     input_shape=(150, 150, 3),  
4     include_top=False,  
5 ) # Do not include the ImageNet classifier at the top.  
6  
7 # Freeze the base_model  
8 base_model.trainable = False
```

```
9
10 # Create new model on top
11 inputs = keras.Input(shape=(150, 150, 3))
12
13 # Pre-trained Xception weights requires that input be scaled
14 # from (0, 255) to a range of (-1., +1.), the rescaling layer
15 # outputs: `(inputs * scale) + offset`
16 scale_layer = keras.layers.Rescaling(scale=1 / 127.5, offset=-1)
17 x = scale_layer(inputs)
18
19 # The base model contains batchnorm layers. We want to keep them in inference mode
20 # when we unfreeze the base model for fine-tuning, so we make sure that the
21 # base_model is running in inference mode here.
22 x = base_model(x, training=False)
23 x = keras.layers.GlobalAveragePooling2D()(x)
24 x = keras.layers.Dropout(0.2)(x) # Regularize with dropout
25 outputs = keras.layers.Dense(1)(x)
26 model = keras.Model(inputs, outputs)
27
28 model.summary(show_trainable=True)
```

📄 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_83683744/83683744 1s 0us/step
Model: "functional"

Layer (type)	Output Shape	Param #	Traina...
input_layer_1 (InputLayer)	(None, 150, 150, 3)	0	-
rescaling (Rescaling)	(None, 150, 150, 3)	0	-
xception (Functional)	(None, 5, 5, 2048)	20,861,480	N
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0	-
dropout (Dropout)	(None, 2048)	0	-
dense (Dense)	(None, 1)	2,049	Y

Total params: 20,863,529 (79.59 MB)
Trainable params: 2,049 (8.00 KB)
Non-trainable params: 20,861,480 (79.58 MB)

```
1 model.compile(
2     optimizer=keras.optimizers.Adam(),
3     loss=keras.losses.BinaryCrossentropy(from_logits=True),
4     metrics=[keras.metrics.BinaryAccuracy()],
5 )
6
7 epochs = 2
8 print("Fitting the top layer of the model")
9 model.fit(train_ds, epochs=epochs, validation_data=validation_ds)
```

📄 Fitting the top layer of the model
Epoch 1/2
146/146 100s 568ms/step - binary_accuracy: 0.8537 - loss: 0.3007 - val_binary_accuracy: 0.9699 - val_loss: 0.1322
Epoch 2/2
146/146 18s 126ms/step - binary_accuracy: 0.9422 - loss: 0.1322 - val_binary_accuracy: 0.9716 - val_loss: 0.1322
<keras.src.callbacks.history.History at 0x7a2f291cbfd0>

```
1 # Unfreeze the base_model. Note that it keeps running in inference mode
2 # since we passed `training=False` when calling it. This means that
3 # the batchnorm layers will not update their batch statistics.
4 # This prevents the batchnorm layers from undoing all the training
5 # we've done so far.
6 base_model.trainable = True
7 model.summary(show_trainable=True)
8
9 model.compile(
10     optimizer=keras.optimizers.Adam(1e-5), # Low learning rate
11     loss=keras.losses.BinaryCrossentropy(from_logits=True),
12     metrics=[keras.metrics.BinaryAccuracy()],
13 )
14
15 epochs = 1
16 print("Fitting the end-to-end model")
17 model.fit(train_ds, epochs=epochs, validation_data=validation_ds)
```



Model: "functional"

Layer (type)	Output Shape	Param #	Traina...
input_layer_1 (InputLayer)	(None , 150, 150, 3)	0	-
rescaling (Rescaling)	(None , 150, 150, 3)	0	-
xception (Functional)	(None , 5, 5, 2048)	20,861,480	Y
global_average_pooling2d (GlobalAveragePooling2D)	(None , 2048)	0	-
dropout (Dropout)	(None , 2048)	0	-
dense (Dense)	(None , 1)	2,049	Y

Total params: 20,867,629 (79.60 MB)
Trainable params: 20,809,001 (79.38 MB)
Non-trainable params: 54,528 (213.00 KB)
Optimizer params: 4,100 (16.02 KB)

Fitting the end-to-end model
146/146 137s 577ms/step - binary_accuracy: 0.8508 - loss: 0.3771 - val_binary_accuracy: 0.9643 - val_loss: 0.1073535904288292
<keras.src.callbacks.history.History at 0x7a2fbec9f040>

- 1 print("Test dataset evaluation")
- 2 model.evaluate(test_ds)



Test dataset evaluation
37/37 4s 106ms/step - binary_accuracy: 0.9614 - loss: 0.0964
[0.1073535904288292, 0.9582974910736084]

- 1 Start coding or [generate](#) with AI.